

NON ROCKET SCIENCE KUBERNETES

Manuel Calvo

OBJECTIVES

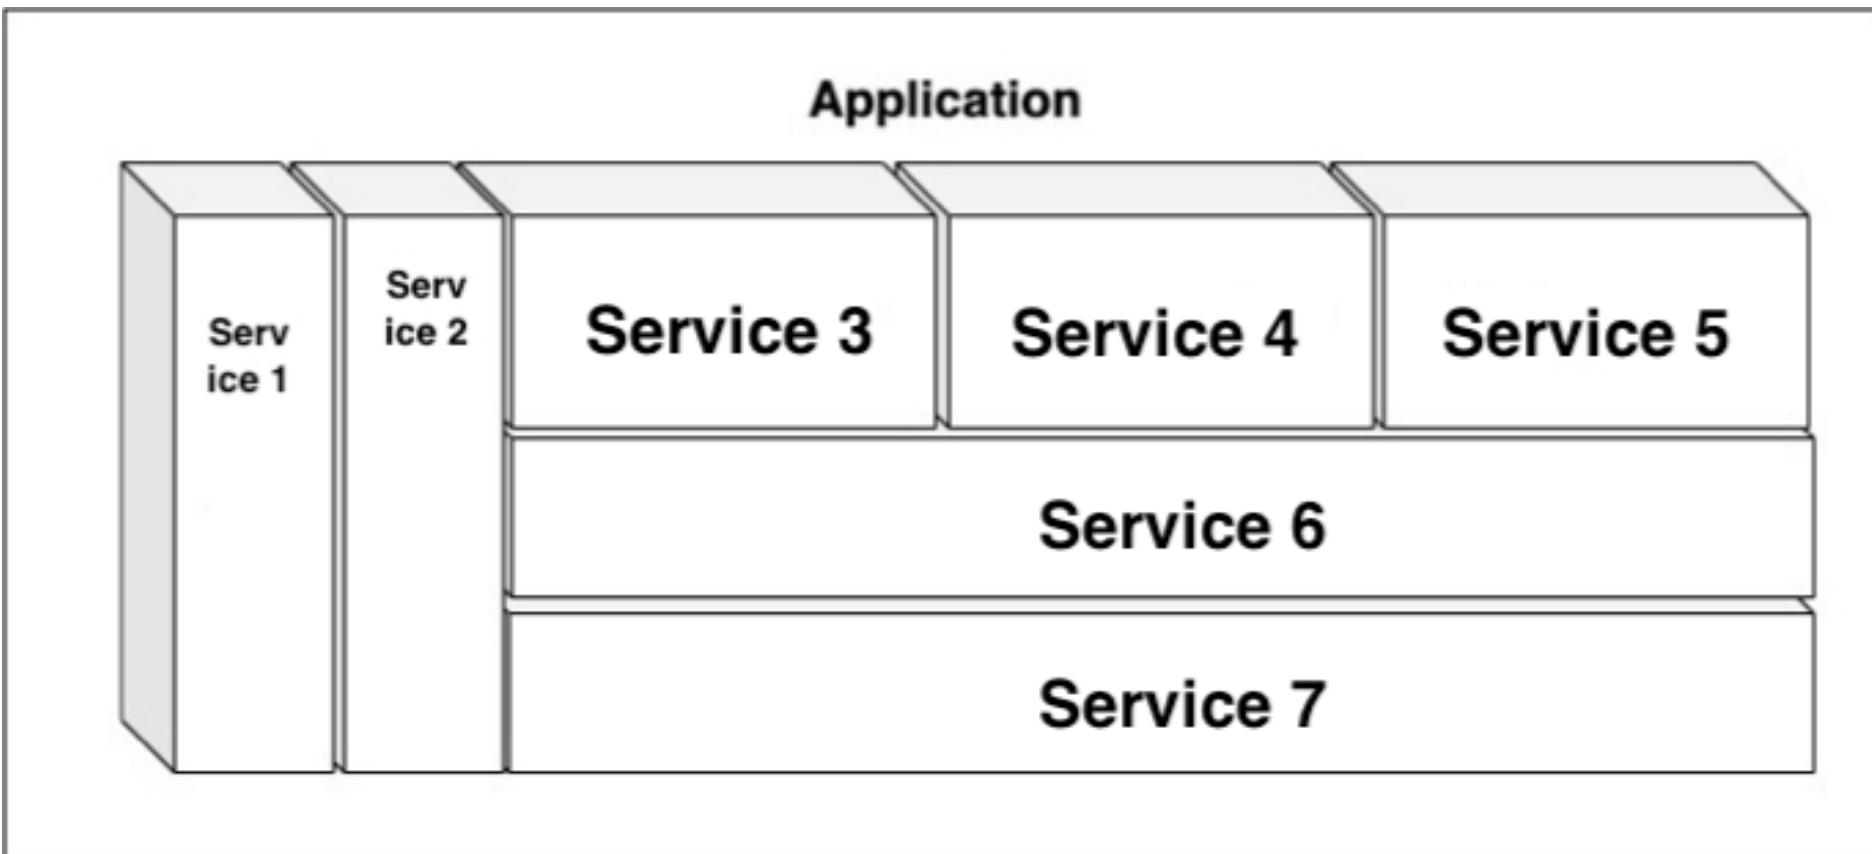
We are not time constrained, we will cover as much as we can in 5 weeks

We are not here just to say I did the course, we are here to help you get started using K8s

Will try to make it hands on as possible

Please interrupt me if you have a question

THE MONOLITH



THE MONOLITH / ADVANTAGES

Only one piece of software: easier testing, trouble shoot, monitoring and management

Need less network traffic

Easier to achieve data consistency

THE MONOLITH / DISADVANTAGES

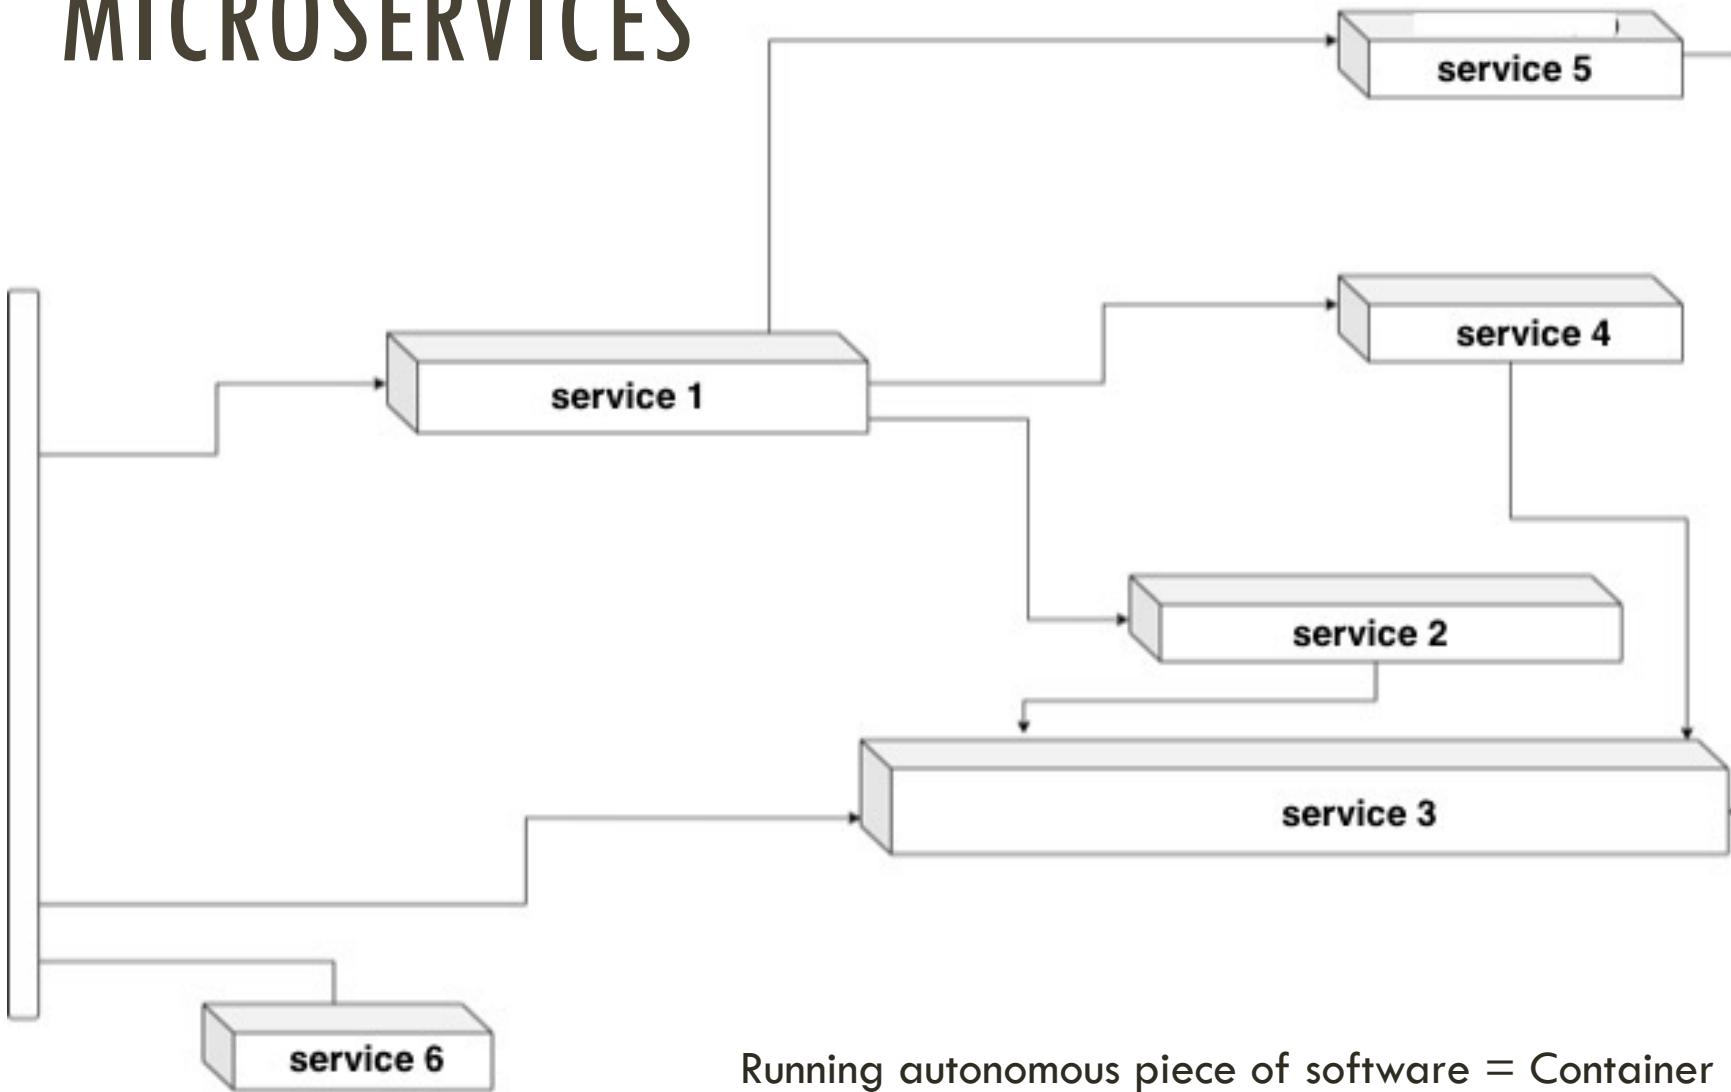
Can't scale individual components

Managing a single piece of code by a large team can become complex and prompt for bottle necks

Harder to mix / incorporate different technologies

Slower time to market

MICROSERVICES



“Microservices are small, autonomous services that work well together”

- Sam Newman

These autonomous pieces of software run in what is known as containers

MICROSERVICES / ADVANTAGES

Can scale components individually

Easier to adopt new technologies

Easier to distribute workload

Faster time to market

MICROSERVICES / DISADVANTAGES

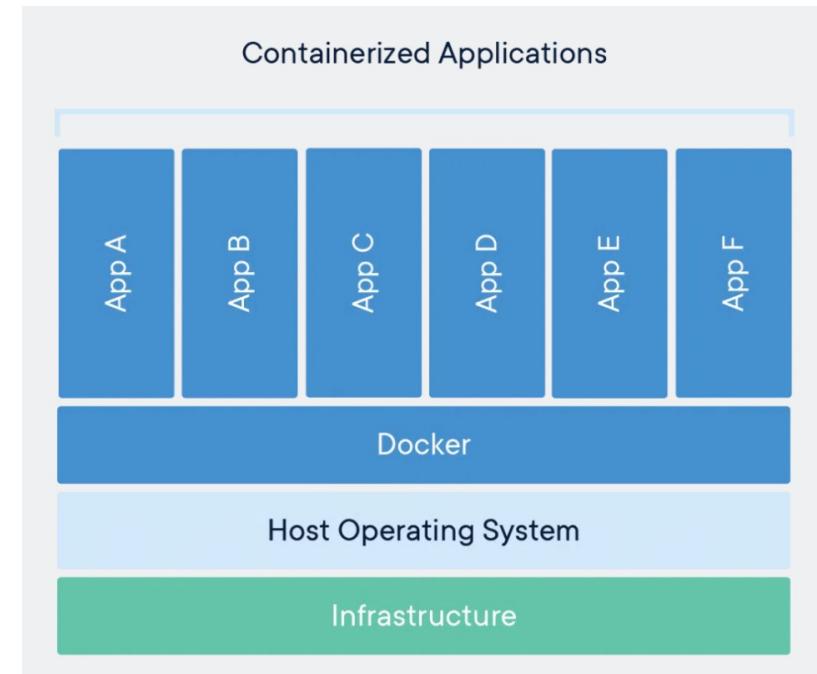
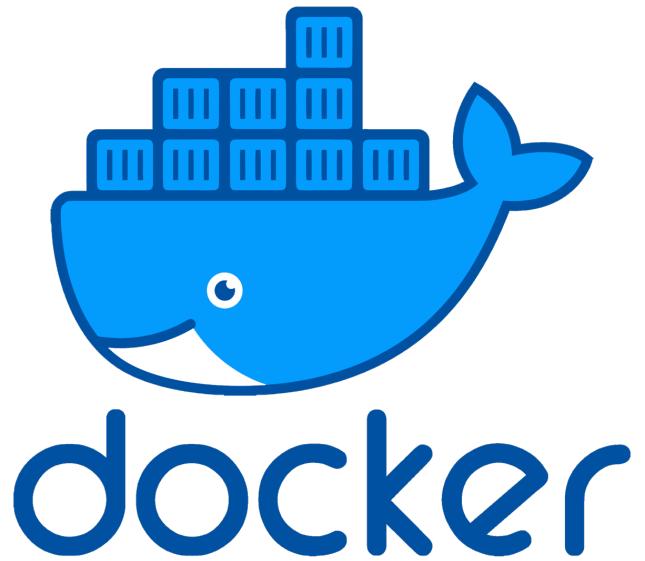
Harder to test, trouble shoot, monitor and manage

Require additional network traffic

Harder to achieve data consistency

More expensive to run

IMAGES AND CONTAINERS



Docker: the most famous tool to work with containers, but not the only option in the market

OCI (Open Container Initiative) = creating open industry standards around container formats and runtimes

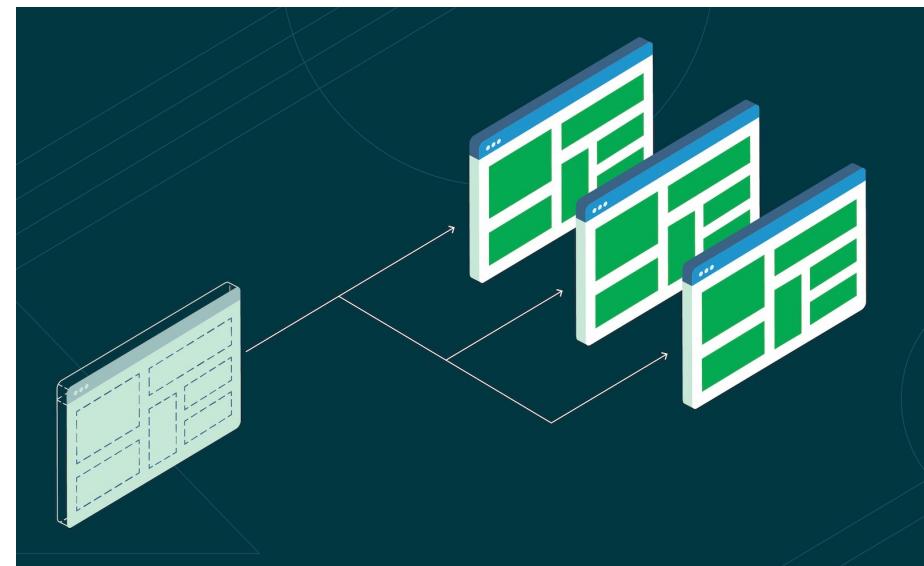
IMAGES AND CONTAINERS

Image: is the actual software we run, it needs to be build in a format containers can manage (OCI standard format)

Container: an actual running instance of the software packaged as an image.

“Think of a container as a shipping container for software — it holds important content like files and programs so that an application can be delivered efficiently from producer to consumer”

- <https://circleci.com/blog/docker-image-vs-container>



Why this picture ?

IMAGES AND CONTAINERS

Containers registries

Some are public some are private

All software that runs in a K8s cluster, runs inside a container and its downloaded from a container registry

<https://hub.docker.com/search>

▪ ie: mysql

- 192.168.99.100:5000/my-image
- myregistry.azurecr.io/marketing/campaign10-18/email-sender:v2
- hello-world:latest

LAB: SEARCH DOCKER HUB

<https://hub.docker.com/search>

Search docker hub and give us an example of a container image that caught your attention

Why that image ?

What version you think you will need ?

What is its image URI ?

WHY K8S ?

Besides just running the software in the containers, K8s adds additional features:

Self healing (ie: if a container goes down or unresponsive, it will automatically restart it for us, will reclaim resources based on priorities)

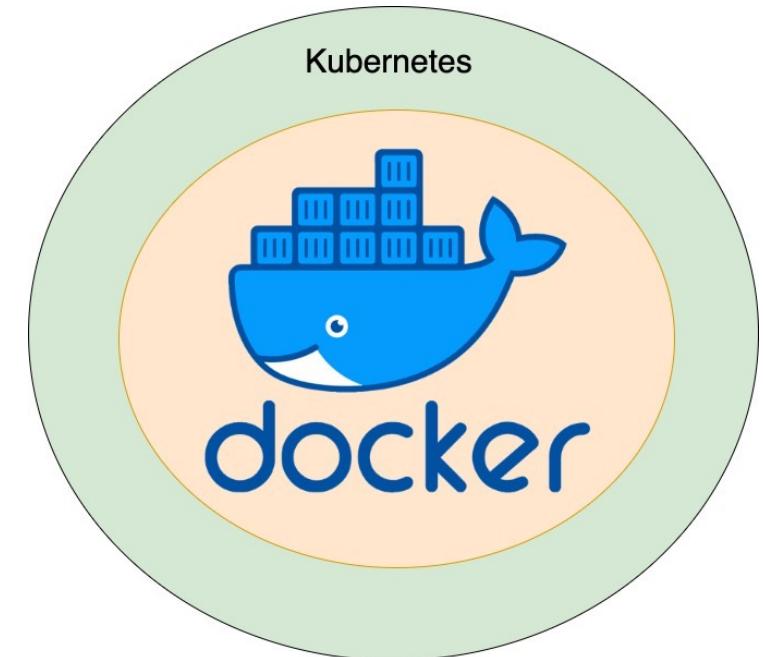
Cluster, multi node (unlike Docker, where a single host manages all the containers, Kubernetes clusters distribute the workload across multiple machines)

Security (defines who and what can be accessed)

Sharing information

Provide ways to organize large deployments (multiple cluster, partition clusters, distribute loads)

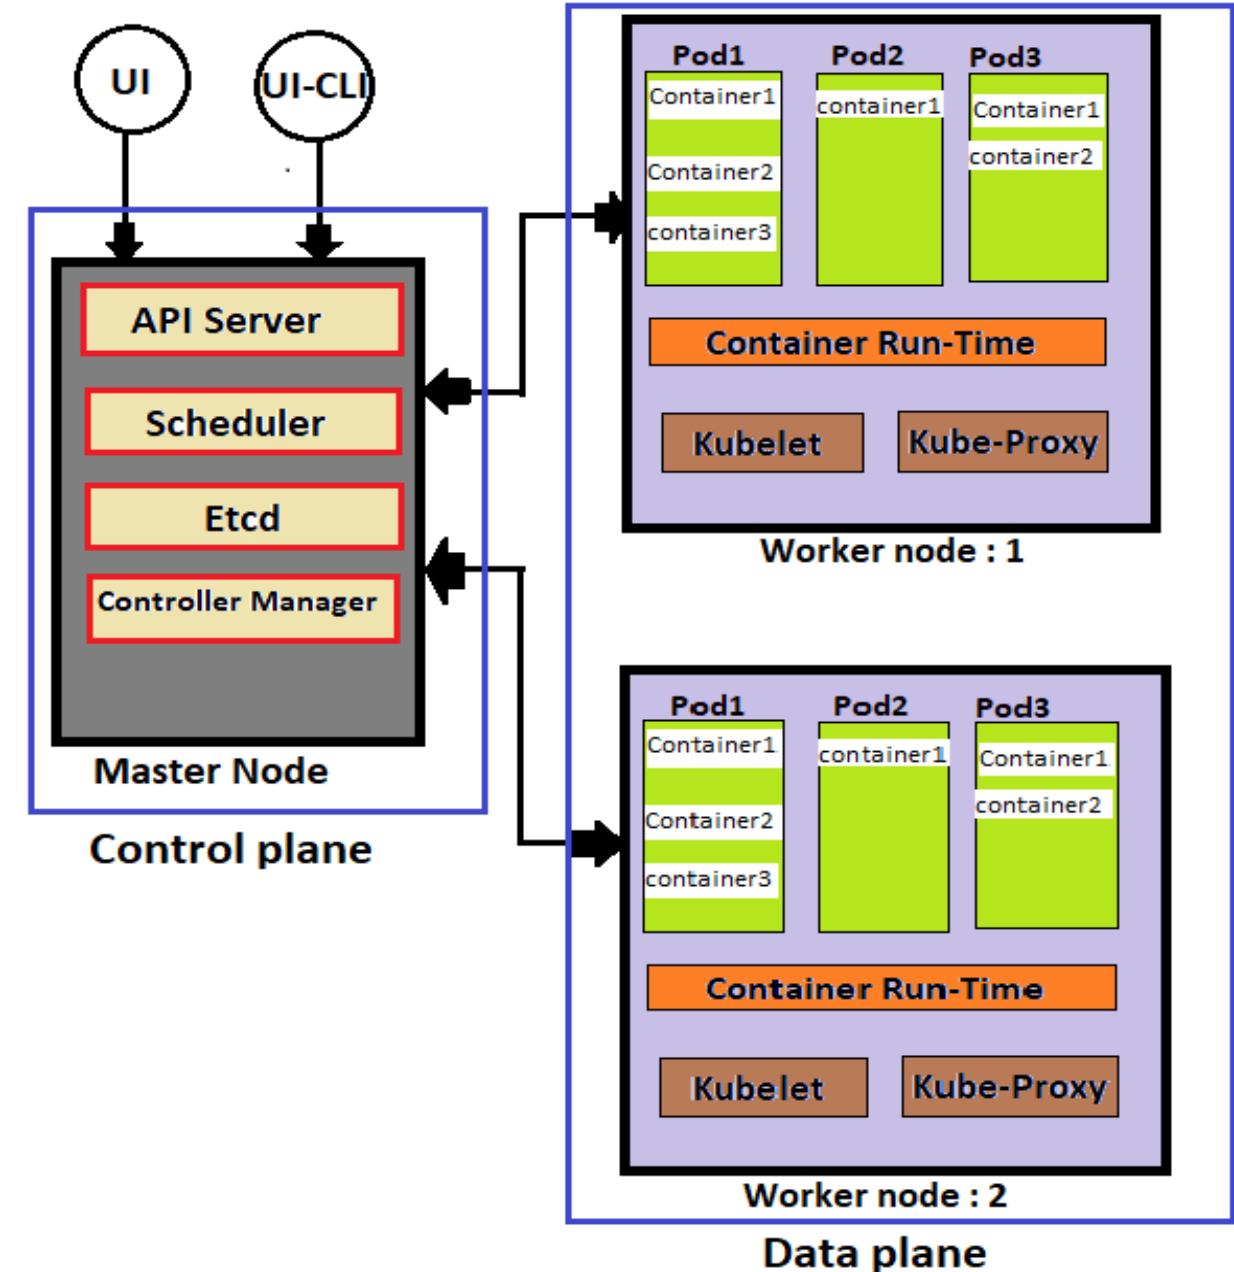
Declarative (you just tell K8s what you want, and it will determine what is the best way to achieve that request)



More than just running containers, K8s
is a container orchestrator

K8S ARCHITECTURE

Kubernetes consists of two main parts: the **control plane (master node)** and the **data plane (worker node)**. Each part consists of multiple components



K8S ARCHITECTURE / CONTROL PLANE

API Server: acts as the central communication hub within the Kubernetes cluster. End users, and other cluster components, talk to the cluster via the API server. When you use *kubectl* you are actually communicating with the API server through HTTP REST APIs. It is the only component that communicates with etcd

Etcd: serves as the backing store for the entire cluster's information. It stores the configuration and state data of the cluster

Scheduler: is responsible for assigning tasks, such as pods to worker nodes based on available resources and constraints. It ensures optimal resource utilization and load balancing across the cluster. It is a controller

Controller Manager: ensures that the cluster remains in the desired configuration at all times. Monitors the state of the cluster and takes corrective actions if anything goes wrong. Some built-in controllers: replicaset, daemonset, cronjob. You can extend kubernetes with custom controllers associated with a custom resource definition (CRDs, Operator pattern)

K8S ARCHITECTURE / WORKER NODE

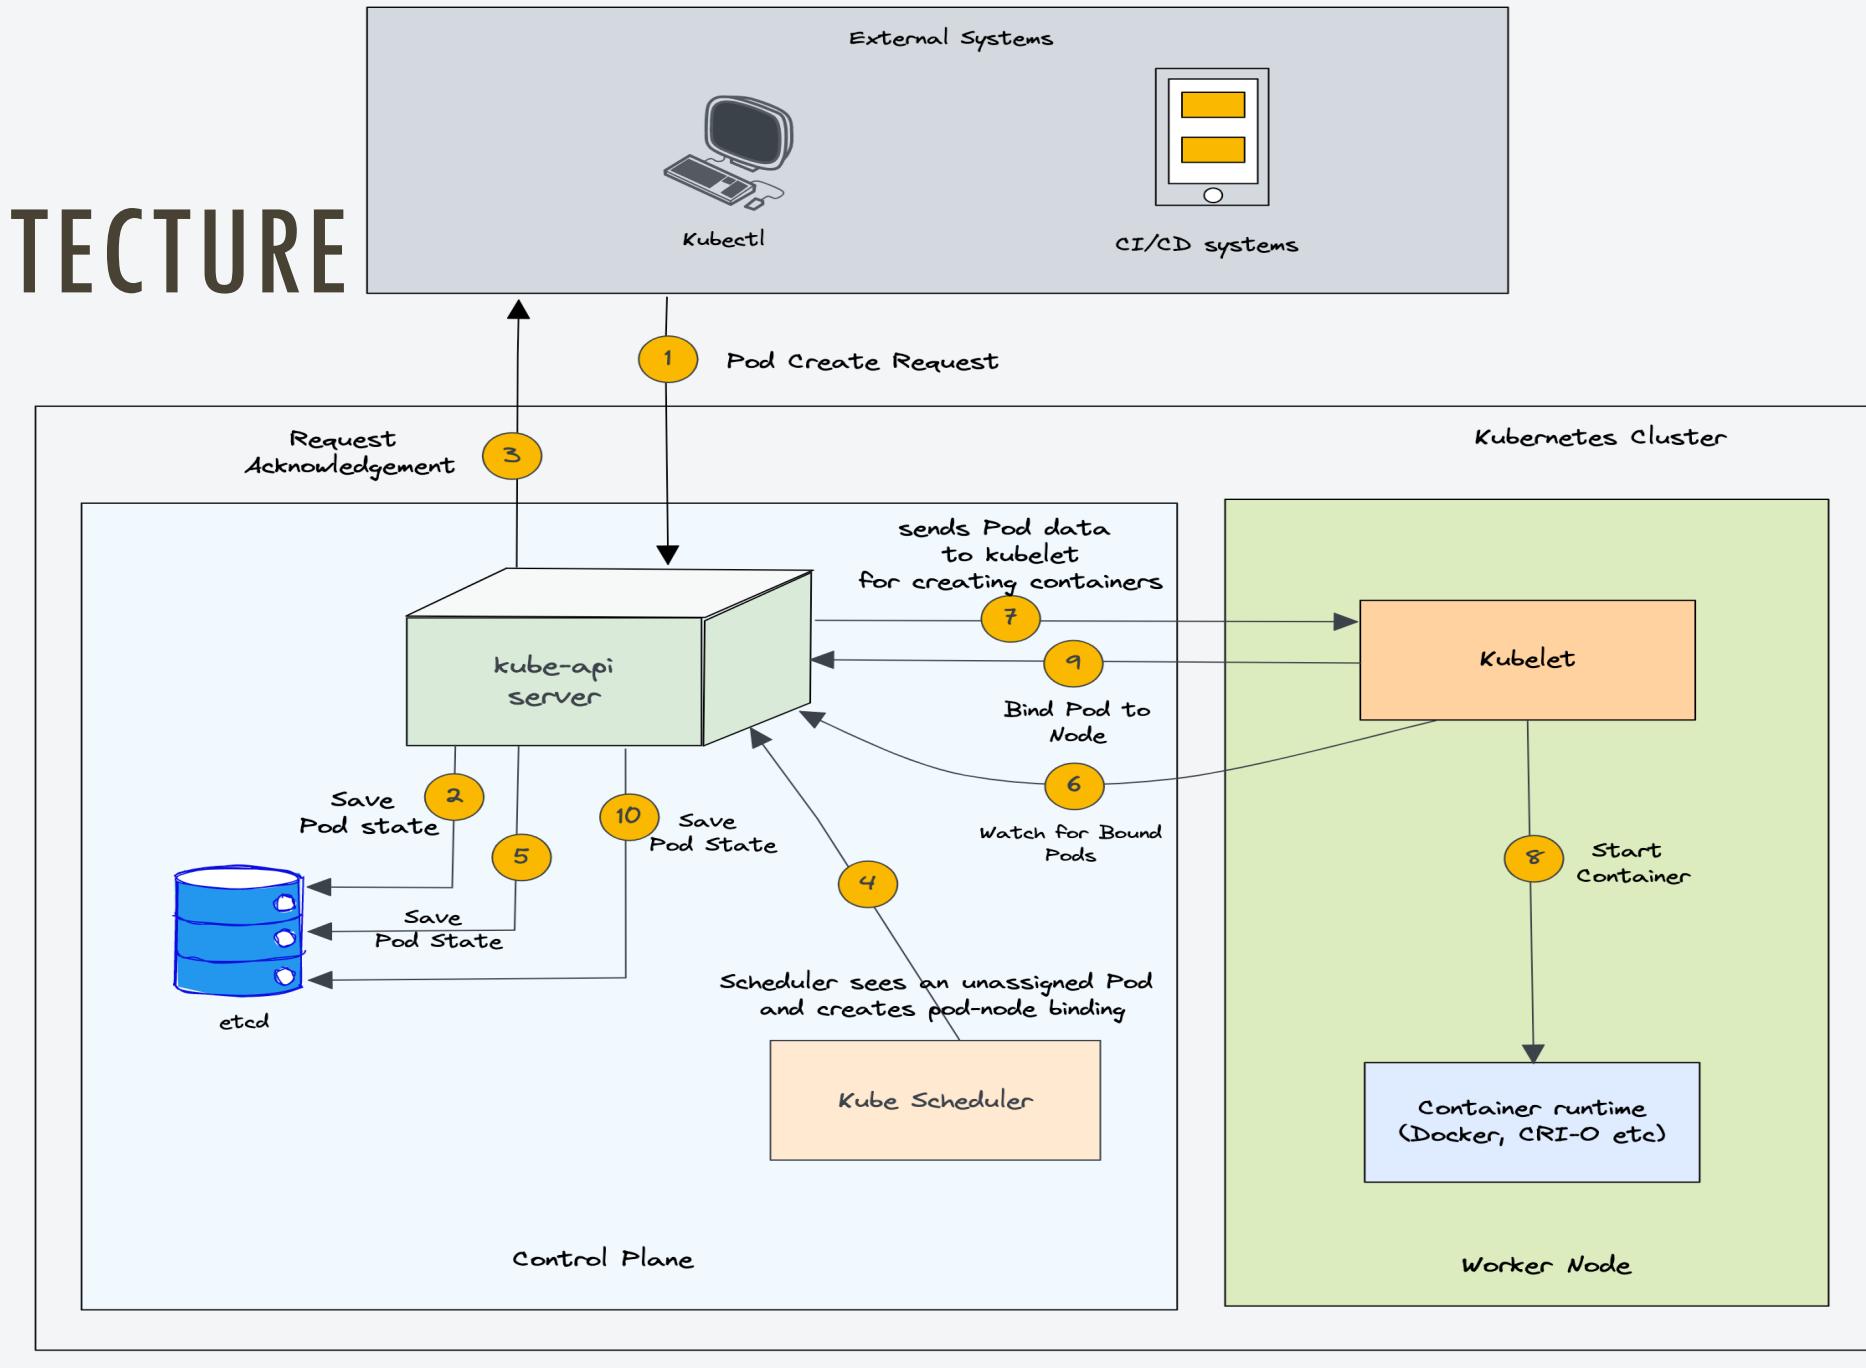
Kubelet: manages the state of individual worker nodes. It ensures that pods are created and running on the node. If a pod fails, the kubelet informs the API server, which then instructs the controller manager to regenerate the container. Mounts volumes. It manages Pods networking through the Container Networking Interface (CNI)

Kube-proxy: handles networking tasks within *services* in the cluster. It generates IP addresses and performs load balancing to enable communication. It utilizes IP tables on the Linux machine to configure networking

Container Runtime: is the software responsible for running and managing containers. It provides the environment in which containers can execute, including managing the lifecycle of containers, resource allocation, and isolation

- **Container Runtime Interface (CRI):** It is a set of APIs that allows Kubernetes to interact with different container runtimes. It allows different container runtimes to be used interchangeably with Kubernetes

K8S ARCHITECTURE



LAB: INSTALL MINIKUBE

Minikube is a lightweight Kubernetes implementation that allows to run a simple cluster on your local machine. There are other alternatives like: k3s, kind

<https://minikube.sigs.k8s.io/docs/start/>

macOS

- brew **install** minikube

Windows

- winget **install** minikube

LAB: INSTALL MINIKUBE

```
minikube config set memory 8g
```

```
minikube config set cpus 4
```

```
minikube config set driver hyperkit (only MacOs)
```

```
minikube config set container-runtime docker
```

```
minikube -p learn start --kubernetes-version=v1.18.6 --nodes=2
```

LAB: KUBECTL / CHECKING INSTALLATION

What is kubectl?

- alias k=kubectl
- nano ~/.zshrc or ~/.bashrc
- source ~/.zshrc

// Gets K8s version number

```
kubectl version --output=json
```

// Gets list of types of resources installed in the cluster

```
kubectl api-resources
```

LAB: KUBECTL / CHECKING INSTALLATION

```
// Gets list of resources in default namespace  
kubectl get all
```

```
// List of resources of a specific type  
kubectl get <resource_type or short_name> -o wide --watch  
▪ kubectl get nodes -o wide  
▪ Kubectl get ev -o wide --watch  
▪ kubectl get po -o wide --watch // List pods. Why po?
```

```
// Gets details about a specific resource  
kubectl describe <resource_type or short_name> <resource_name>  
▪ kubectl describe node learn
```

LAB: KUBECTL / CHECKING INSTALLATION

```
// Gets logs  
kubectl logs <pod_name or service_name>  
kubectl logs --tail=20 <pod_name or service_name> // last 20 entries  
kubectl logs -f <pod_name or service_name> // monitor for new entries
```

Difference between logs and events

LAB: HELLO WORLD

// This creates a **pod** with the name *first* and passes some env variables

// That pod is running your image in a container

```
kubectl run first --image=cloudnative/demo:hello --env="TEXT=Hi" --env="OTHER=value2"
```

<https://hub.docker.com/search>

Tell us:

- Is it running?
- In which node it runs?
- What is its IP?
- Can you ping it?

<https://github.com/shiriuki/k8s-training/blob/main/kubectl-commands.yaml>

LAB: HELLO WORLD

Why are you not able to ping your pod?

How to ping your pod:

- `m ssh -p learn -n learn-m02` // Connects through ssh to the “node” that is running the pod.
- `ping 10.224.1.4`
- `curl http://10.244.1.4:8080` // If it exposes a web port

A simpler way, open a temporary channel through the kubectl to access the pod’s containers ports (advantages don’t need to the pod’s IP or machine where it is running):

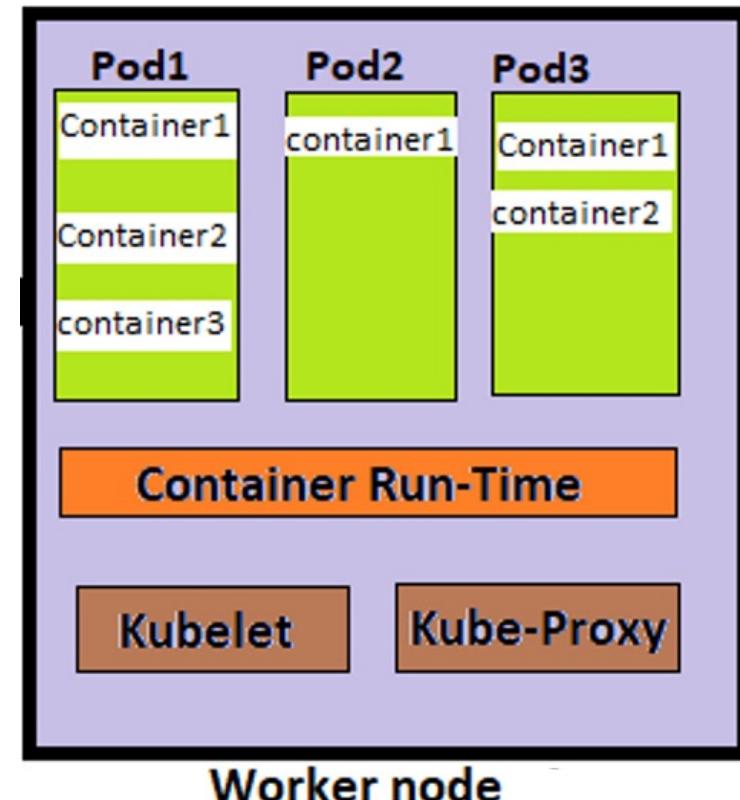
```
kubectl port-forward <pod_name> <local_port>:<pod_container_port>
kubectl port-forward first 7777:8080
curl http://localhost:7777
```

PODS

Pod is most granular unit of execution in k8s. Anything that you want to run in K8s, runs as an image in a container, and that container belongs to a pod. You can't just run a container (as in Docker), the container needs to be part of a pod

Each pod has a unique IP address and the port space is shared by all the containers in that pod. This means that different containers inside a pod can communicate with each other using their corresponding ports on localhost. They also share volumes.

```
// connects to the pod's container  
kubectl exec -it pod_name -c container_name bash
```



PODS

```
// Generates the yaml representation of a K8s resource  
kubectl get pod first -o yaml > firstPod.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: first  
spec:  
  containers:  
    - env:  
        - name: TEXT  
          value: Hi  
        - name: OTHER  
          value: value2  
      image: cloudnativd/demo:hello  
      imagePullPolicy: IfNotPresent  
      name: first
```

LAB: PODS

`kubectl explain pod // Gives information about a K8s resource and its fields`

`kubectl explain pod.spec.containers.imagePullPolicy // Gives info about a field`

`kubectl apply -f secondPod.yaml // deploys a K8s resource described in yaml`

Deploy a new Pod:

- **name:** second
- **image:** wardviaene/http-echo
- should display a customized message (how ?)
- should always download the image (how ?)
- What is your approach to do it?
- Run **kubectl get pods -o wide -watch** in one tab, and the command to deploy the new pod in another tab

<https://github.com/shiriuki/k8s-training/blob/main/kubectl-commands.yaml>

POD PATTERNS

Remember a Pod can run 1 or more containers

```
kind: Pod
apiVersion: v1
metadata:
  name: multi
spec:
  containers:
    - name: container-1
      image: nginx
    - name: container-2
      image: ubuntu
      command:
        - /bin/bash
        - -c
        - while :; do echo '.'; sleep 5; done
```

<https://betterprogramming.pub/understanding-kubernetes-multi-container-pod-patterns-577f74690aee>

POD PATTERNS / INIT CONTAINER

Init container

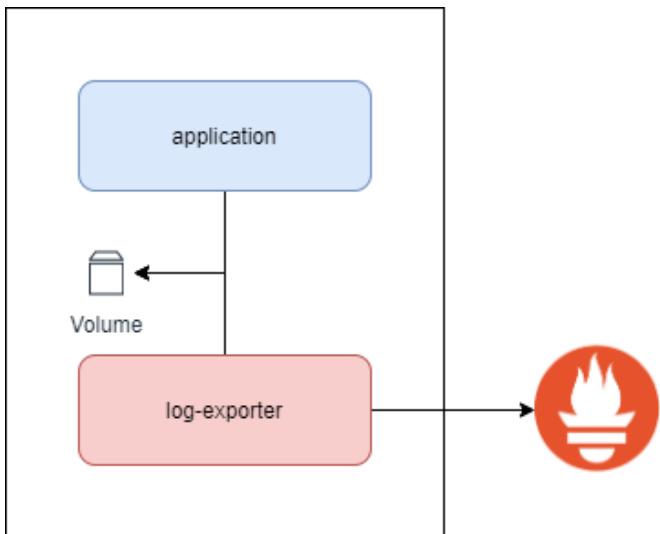
The init container must run successfully before the regular containers can run. If that doesn't happen K8s won't proceed to execute the regular containers

```
kind: Pod
apiVersion: v1
metadata:
  name: example
spec:
  initContainers:
  - name: create-db
    image: postgres:12
    command:
      - /bin/bash
      - -c
      - psql -h postgresHost -U postgres -c "CREATE DATABASE MYDB"
  containers:
  - name: activity-sender
    image: gcr.io/radix/activity-sender
```

POD PATTERNS / SIDECAR

Sidecar

Is used to enhance or extend the existing functionality of the main container. Your container works perfectly well without the sidecar, but with it, it can perform some extra functions.

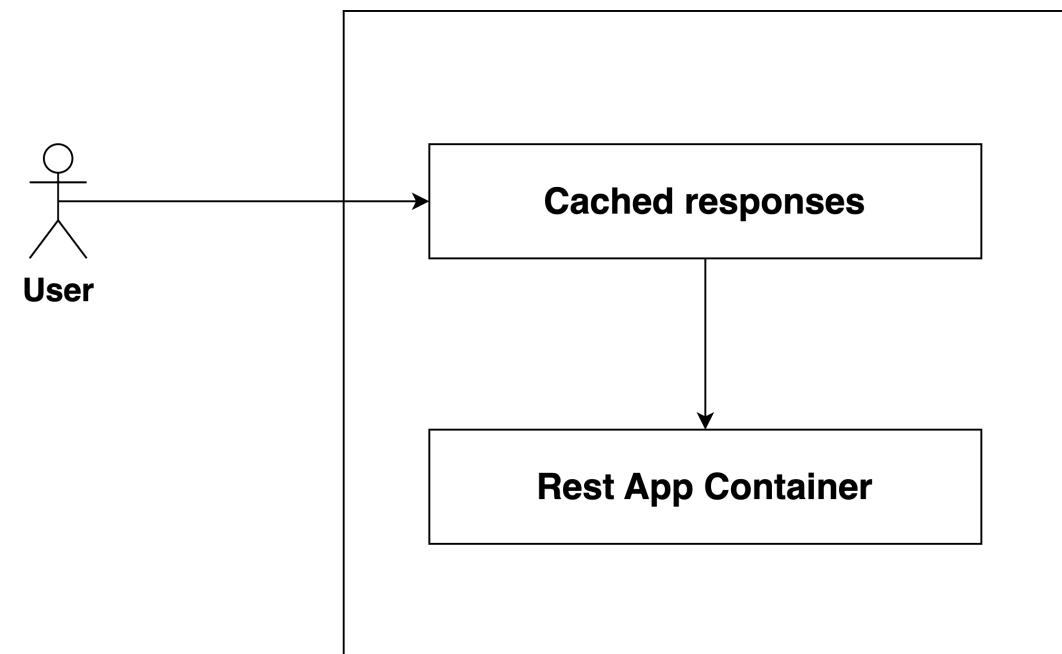


```
apiVersion: v1
kind: Pod
metadata:
  name: sidecar-pod
spec:
  volumes:
  - name: logs
    emptyDir: {}
  containers:
  - name: app-container
    image: alpine
    command: ["/bin/sh"]
    args: ["-c", "while true; do date >> /var/log/app.log; sleep 2;done"]
    volumeMounts:
    - name: logs
      mountPath: /var/log
  - name: log-exporter-sidecar
    image: nginx
    ports:
    - containerPort: 80
    volumeMounts:
    - name: logs
      mountPath: /usr/share/nginx/html
```

POD PATTERNS / AMBASSADOR

Ambassador

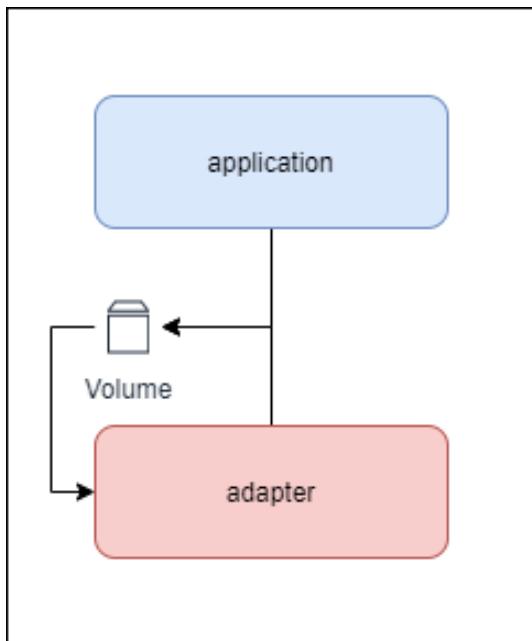
A proxy to the external world. We create another container that can act as a proxy to the main container. For example a cache of responses of the main container



POD PATTERNS / ADAPTER

Adapter

The adapter pattern helps you standardize something. For example, transforming the logs of a legacy app to a newer standard format.



```
apiVersion: v1
kind: Pod
metadata:
  name: adapter-pod
  labels:
    app: adapter-app
spec:
  volumes:
  - name: logs
    emptyDir: {}
  containers:
  - name: app-container
    image: alpine
    command: ["/bin/sh"]
    args: ["-c", "while true; do date >> /var/log/app.log; sleep 2;done"]
    volumeMounts:
    - name: logs
      mountPath: /var/log
  - name: log-adapter
    image: alpine
    command: ["/bin/sh"]
    args: ["-c", "tail -f /var/log/app.log|sed -e 's/^Date /' > /var/log/out.log"]
    volumeMounts:
    - name: logs
      mountPath: /var/log
```

LABELS

With them you can categorize your K8s resources, so later you can filter them out. They are key/value pairs

Ie: **google.com/customer=ibm**, **google.com/customer=common**

// Gets all resources and their current labels

```
kubectl get all --show-labels -A
```

// Filters resources by label

```
kubectl get all -A -l tier=control-plane
```

```
kubectl get pods -A -l google.com/customer=ibm, google.com/bu!=finance
```

LABELS

Valid key has 2 segments:

optional prefix and name, separated by a slash (/). If specified,
the prefix must be a DNS not longer than 253 characters

ie: google.com/app

Rules for valid label value and name segment:

no more than 63 characters

must begin and end with an alphanumeric character ([a-z0-9A-Z])

could contain dashes (-), underscores (_), dots (.)

```
// Adds labels to an existing K8s resource
kubectl label <resource_type> <resource_id> google.com/nodetype=db
kubectl label node learn-m02 google.com/nodetype=db -overwrite
kubectl label all -l k8s-app=kube-dns -A mylabel=value
```

```
apiVersion: v1
kind: Pod
metadata:
  name: declarative-nginx
  labels:
    radixiot.com/app: nginx
spec:
  containers:
  - name: nginx
    image: nginx
```

ANNOTATIONS

Used by third party apps to attach metadata to K8 resources. K8s doesn't care or know about that metadata or how it is used. The keys & values used in annotations have no restrictions

```
kind: Pod
apiVersion: v1
metadata:
  name: pod-with-annotation
  annotations:
    velero.io/dontBackup: "true"
spec:
  containers:
  - name: first-container
    image: nginx
```

```
// Adds annotations to an existing K8s resource
kubectl annotate pod nginx \
  pre.hook.backup.velero.io/command='["/sbin/fsfreeze", "--freeze", "/var/log/nginx"]' \
  post.hook.backup.velero.io/command='["/sbin/fsfreeze", "--unfreeze", "/var/log/nginx"]'

// To remove annotation
kubectl annotate pod nginx post.hook.backup.velero.io/command-
```

SERVICES

A service associates a set of pods with an abstract service name

The need for services arises from the fact that pods are short lived and can be replaced at any time. This means that pods that need to communicate with another pod cannot rely on the IP address of the underlying single pod. Instead, they connect to the service name, which relays them to a relevant, currently-running pod

Uses labels and selectors to match the pods it represents

It will distribute the load between those pods

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
```

SERVICES / TYPES

ClusterIP

It receives a cluster-internal IP address, making its pods only accessible from within the cluster

Internal K8s DNS creates an entry that allows to resolve:

- <http://my-cip-service.default:7777>
- <http://my-cip-service.default.svc.cluster.local:7777>
- <http://my-cip-service:7777>
- k exec -it pod1 sh
- nslookup my-cip-service 8.8.8.8

TIPS:

kubectl port-forward service/myservice 80:7777 doesn't distribute the load, also a gRPC channel doesn't distribute the load
you can't ping a service IP

```
apiVersion: v1
kind: Service
metadata:
  name: my-cip-service
spec:
  selector:
    app: msg
  type: ClusterIP
  ports:
  - protocol: TCP
    port: 7777
    targetPort: 8080
```

SERVICES / TYPES

Behind the scenes K8s creates (and automatically adjust) a list of pods that match the service's selector

kubectl get endpoints

```
[manuelcalvo [Downloads] # k get endpoints
NAME           ENDPOINTS             AGE
kubernetes     192.168.205.77:8443   11d
my-cip-service 10.244.0.3:8080,10.244.1.4:8080,10.244.1.5:8080   79m
```

TIP:

You can create a service without a selector, and manually define its endpoints. This can be used to access services outside the K8s cluster that can't be discovered through DNS

```
apiVersion: v1
kind: Service
metadata:
  name: minio-service
  namespace: velero
spec:
  type: ClusterIP
  ports:
  - protocol: TCP
    port: 9000
    targetPort: 9000
---
apiVersion: v1
kind: Endpoints
metadata:
  name: minio-service
  namespace: velero
subsets:
- addresses:
  - ip: 192.168.2.193
    ports:
    - port: 9000
```

SERVICES / TYPES

NodePort

Builds on top of the ClusterIP service, exposing it to a port accessible from outside the cluster

The kube-proxy component on each node is responsible for listening on the node's external ports and forwarding client traffic from the NodePort to the ClusterIP

By default, all nodes in the cluster listen on the service's NodePort, even if they are not running a pod that matches the service selector

The external port must be in the range: 30000-32767

```
apiVersion: v1
kind: Service
metadata:
  name: my-np-service
spec:
  selector:
    app: msg
  type: NodePort
  ports:
  - protocol: TCP
    nodePort: 32000
    port: 7777
    targetPort: 8080
```

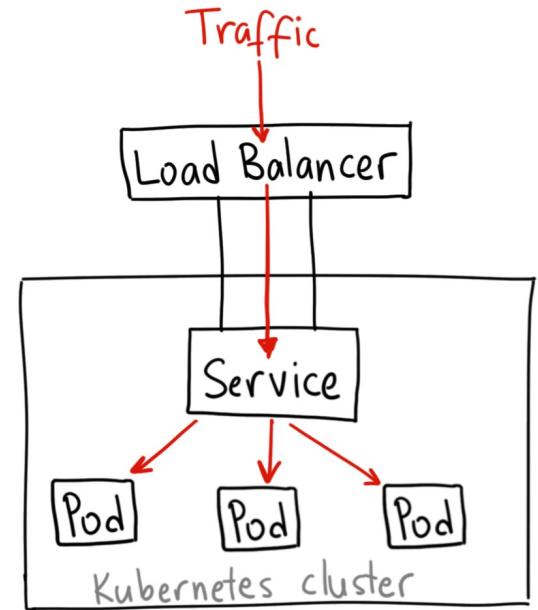
SERVICES / TYPES

LoadBalancer

Builds on top of the NodePort service

To be used in cloud environments for them to provide a load balancer that will forward requests to the K8s cluster

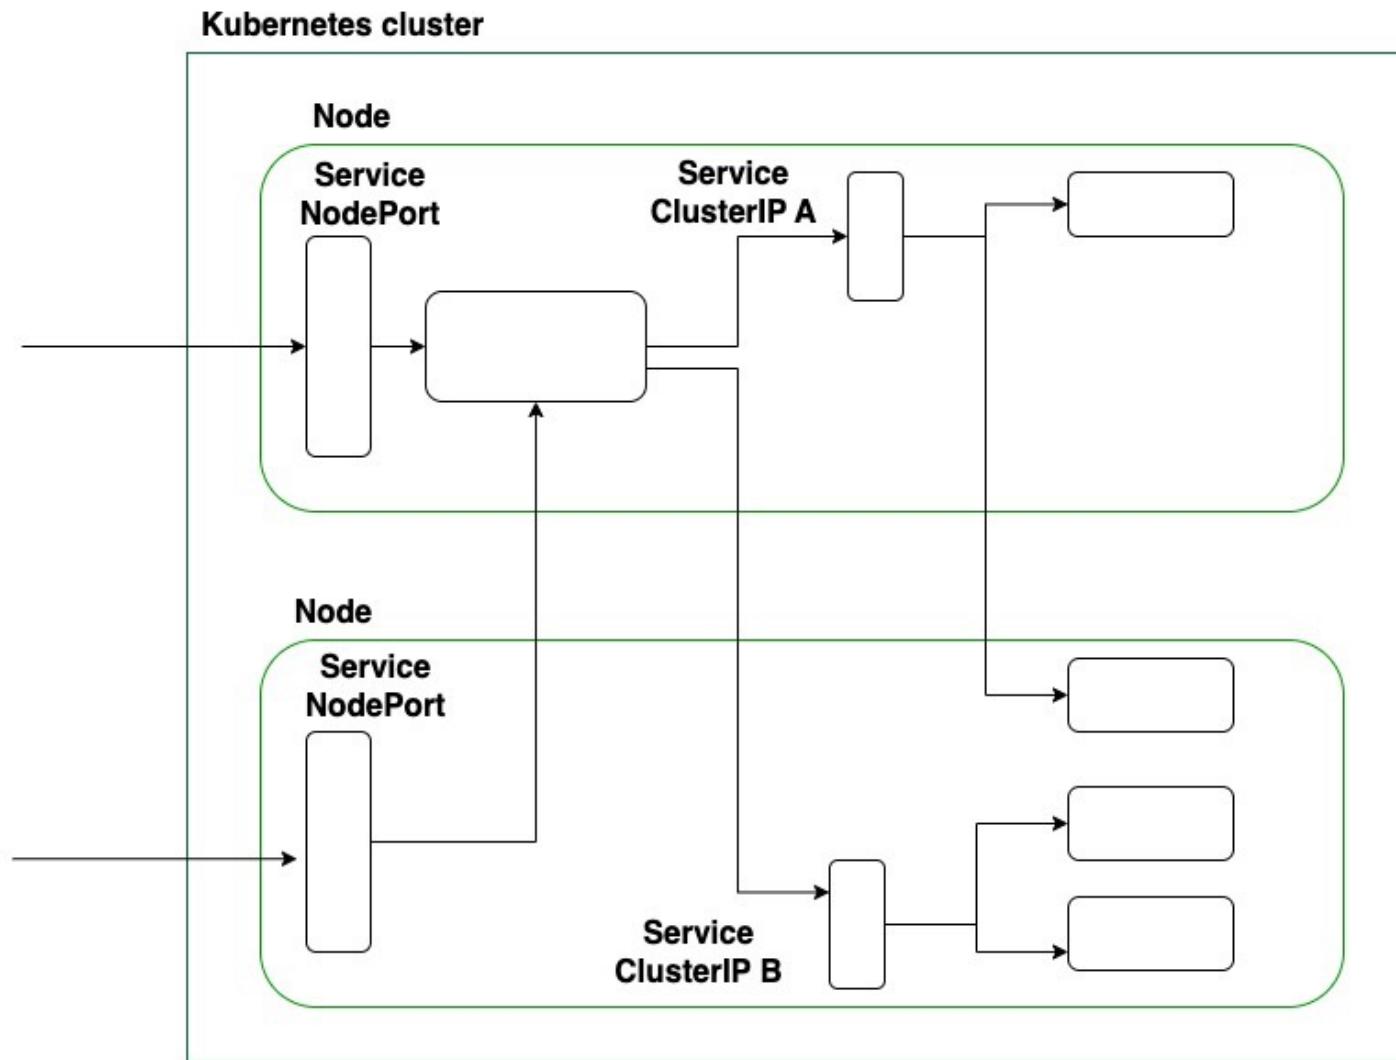
```
apiVersion: v1
kind: Service
metadata:
  name: my-lb-service
spec:
  selector:
    app: msg
  type: LoadBalancer
  ports:
    - protocol: TCP
      nodePort: 30000
      port: 7777
      targetPort: 8080
```



```
$ kubectl get svc nginx-service
NAME        TYPE      CLUSTER-IP     EXTERNAL-IP           PORT(S)        AGE
nginx-service  LoadBalancer  172.20.54.138  ac8415de24f6c4db9b5019f789792e45-443260761.us-east-2.elb.amazonaws.com  80:30968/TCP  21h
```

```
$ curl ac8415de24f6c4db9b5019f789792e45-443260761.us-east-2.elb.amazonaws.com
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

SERVICES / EXAMPLES



LAB: SERVICES

Create 2 pod instances using this image: wardviaene/http-echo

Each pod should display a different message

Then create a service of type NodePort that redirects requests to those pods

Probe load balancing is happening

<https://github.com/shiriuki/k8s-training/blob/main/kubectl-commands.yaml>

<https://github.com/shiriuki/k8s-training/blob/main/service-cip.yaml>

<https://github.com/shiriuki/k8s-training/blob/main/service-np.yaml>

NAMESPACES

Provide a way to partition a cluster

By default, a Kubernetes cluster has a default namespace that is used to deploy resources if no other namespace is specified

They help to improve resource management and enhance security

```
// Create a namespace
kubectl create namespace infra

// Gets the existing namespaces
kubectl get ns

// Gets list of all resources in all namespaces
kubectl get all -A

// Get list of resources of a type in a specific namespace
kubectl get <resource_type> -n <namespace>
kubectl get pod -n kube-system
```

NAMESPACES

```
apiVersion: v1
kind: Namespace
metadata:
  name: infra
```

```
apiVersion: v1
kind: Pod
metadata:
  name: pod1
  namespace: infra
...
...
```

LAB: NAMESPACES

In a **single** yaml file (use --- to indicate the start of a new k8s resource)

- Create a namespace with name: infra
- Deploy Postgres behind a NodePort service in namespace infra
- https://hub.docker.com/_/postgres#:~:text=have%20found%20useful.-,Environment%20Variables,-The%20PostgreSQL%20image
- Postgres default container port 5432

Connect to your Postgres using pgAdmin

- <https://www.pgadmin.org/download/>

List all resources available in all namespaces

```
https://github.com/shiriuki/k8s-training/blob/main/kubectl-commands.yaml
https://github.com/shiriuki/k8s-training/blob/main/service-np.yaml
https://github.com/shiriuki/k8s-training/blob/main/namespace.yaml
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: other
---
apiVersion: v1
kind: Pod
metadata:
  name: pod1
  namespace: other
...
---
apiVersion: v1
kind: Service
metadata:
  name: myservice
  namespace: other
spec:
  selector:
    mylabel: value
...
```

CONFIG MAPS

They provide a way to store data in key-value pairs separately from the application. They can be consumed by pods as environment variables or files in a volume

Should not be used to store confidential or sensitive data because they are stored as plain text

Its data can be shared between pods

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-config-map
data:
  currency: USD
  headerColor: black
  welcomeMsg: Hi there
```

```
apiVersion: v1
kind: Pod
metadata:
  name: podm
spec:
  containers:
    - name: echo
      image: wardviaene/http-echo
      env:
        - name: TEXT
          valueFrom:
            configMapKeyRef:
              name: my-config-map
              key: welcomeMsg
```

SECRETS

Should be used to store sensitive information like passwords, keys, and tokens in **base64 format**

They can be consumed by pods as environment variables or files in a volume

```
echo -n "thepassw" | base64  
echo -n "PHN0cmluZz4=" | base64 -d  
k get secret my-secret -o yaml
```

```
apiVersion: v1  
kind: Secret  
metadata:  
  name: my-secret  
data:  
  <key>: <base64_encoded_value>  
  mongoPassw: dGhlcGFzcw==
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: podm  
spec:  
  containers:  
  - name: db  
    image: mongo  
    env:  
    - name: MONGO_INITDB_ROOT_PASSWORD  
      valueFrom:  
        secretRef:  
          name: my-secret  
          key: mongoPassw
```

LAB: SECRETS

Modify your postgres installation (previous lab) so its password comes from a secret

<https://github.com/shiriuki/k8s-training/blob/main/kubectl-commands.yaml>

<https://github.com/shiriuki/k8s-training/blob/main/secret.yml>

VOLUMES

A volume is a directory containing data, which can be accessed by containers in a pod. The location of the directory, the storage media that supports it, and its contents, depend on the specific type of volume being used

Processes running within containers in a pod see a file system view composed of:

- A root file system that matches the content of the container image
- Volumes mounted on the container (if defined). Each volume mounts on a specific path within the container file system

They are defined within the context of a pod, meaning that you can't create a volume by itself. This makes volumes unmanageable and their lifecycle bound directly to the pod. This means that deleting the pod will result in either deleting the volume or liberating it depending on the volume type used.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-config-map
data:
  currency: USD
  headerColor: black
  welcomeMsg: Hi there
---
apiVersion: v1
kind: Pod
metadata:
  name: podv
spec:
  containers:
  - name: echo
    image: wardviaene/http-echo
    volumeMounts:
    - name: cm
      mountPath: /cm
  volumes:
  - name: cm
    configMap:
      name: my-config-map
```

VOLUMES / EMPTYDIR

A temporary directory of the **node**, that shares a pod's lifetime. When the pod is destroyed, it will destroy the shared volume and all its contents

```
apiVersion: v1
kind: Pod
metadata:
  name: sharevol
spec:
  containers:
    - name: c1
      image: centos:7
      command:
        - "bin/bash"
        - "-c"
        - "sleep 10000"
      volumeMounts:
        - name: xchange
          mountPath: "/tmp/xchange"
    - name: c2
      image: centos:7
      command:
        - "bin/bash"
        - "-c"
        - "sleep 10000"
      volumeMounts:
        - name: xchange
          mountPath: "/tmp/data"
  volumes:
    - name: xchange
      emptyDir: {}
```

VOLUMES / HOSTPATH

A specific directory of the **node** is shared

Because it is a specific directory of the node it won't be destroyed when the pod is destroyed

The user running the container should have permissions to access that directory

```
apiVersion: v1
kind: Pod
metadata:
  name: sharevol
spec:
  containers:
    - name: c1
      image: centos:7
      command:
        - "bin/bash"
        - "-c"
        - "sleep 10000"
      volumeMounts:
        - name: xchange
          mountPath: "/tmp/xchange"
    - name: c2
      image: centos:7
      command:
        - "bin/bash"
        - "-c"
        - "sleep 10000"
      volumeMounts:
        - name: xchange
          mountPath: "/tmp/data"
  volumes:
    - name: xchange
      hostPath:
        path: /manuel/data
```

VOLUMES / CONFIGMAP / SECRET

The contents of config maps or secrets are mounted as virtual directory. Where each key is a file, and the value of the key is the content of the corresponding file

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-config-map
data:
  currency: USD
  headerColor: black
  welcomeMsg: Hi there
---
apiVersion: v1
kind: Pod
metadata:
  name: podv
spec:
  containers:
    - name: echo
      image: wardviaene/http-echo
      volumeMounts:
        - name: cm
          mountPath: /cm
  volumes:
    - name: cm
      configMap:
        name: my-config-map
      volumes:
        - name: <volume_name>
          secret:
            secretName: <secret_name>
```

PERSISTENT VOLUMES

They are completely independent from the pod they are attached to and even independent from the host where it's running when it's not a local persistent volume

Persistent Volumes (PV) decouples storage from the Pod and makes them a first class citizen in kubernetes with their own separate manageable object and abstraction layer

PVs allow data sharing between pods by mounting the same PV to various pods

PVs will remain even after the pod is deleted or replaced

Can be provisioned in two ways:

- manually by an administrator
- dynamically using storage classes with a configured provisioner that specifies which volume plugin to be used for provisioning

PV / MANUAL PROVISIONING

An administrator has access to the underlying storage provider and have knowledge of the disk configuration. For this example in GKE, will create a persistent volume using the GCEPersistentDisk plugin

accessModes:

- ReadWriteOnce — the volume can be mounted as read-write by a single node
- ReadWriteOnly — the volume can be mounted as read-only by many nodes
- ReadWriteMany — the volume can be mounted as read-write by many nodes
- ReadWriteOncePod — the volume can be mounted as read-write by single pod

persistentVolumeReclaimPolicy:

- Retain — manual reclamation
- Delete — associated storage asset such as AWS EBS or GCE PD volume is deleted
- Recycle — basic scrub (`rm -rf /thevolume/*`)

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mongodb-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
    - ReadOnlyMany
  volumeMode: Filesystem
  storageClassName: manual
  persistentVolumeReclaimPolicy: Retain
  gcePersistentDisk:
    pdName: mongodb
    fsType: ext4
```

PV / MANUAL PROVISIONING

Through a PersistentVolumeClaim (PVC) users can request storage by specifying the storage class, accessMode and the size of the volume needed

Binding criteria:

- accessMode must be matched between the PV and PVC.
- The requested storage by a PVC must be equal to or less than the storage capacity of a PV. An equal match will always get priority. But if there is no exact match, then PVC will claim a PV with a larger storage capacity than what was demanded by PVC

Should be " "
for manual
provisioning

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mongodb-pvc
spec:
  resources:
    requests:
      storage: 1Gi
  accessModes:
  - ReadWriteOnce
  storageClassName: "fast"
---
apiVersion: v1
kind: Pod
metadata:
  name: mongodb
spec:
  containers:
  - name: mongo
    image: mongo
    volumeMounts:
    - name: mongodb-data
      mountPath: /data/db
  volumes:
  - name: mongodb-data
    persistentVolumeClaim:
      claimName: mongodb-pvc
```

PV / DYNAMIC PROVISIONING

An K8S object called storage class (“StorageClass”) dynamically creates persistent volumes when requested using a PersistentVolumeClaim

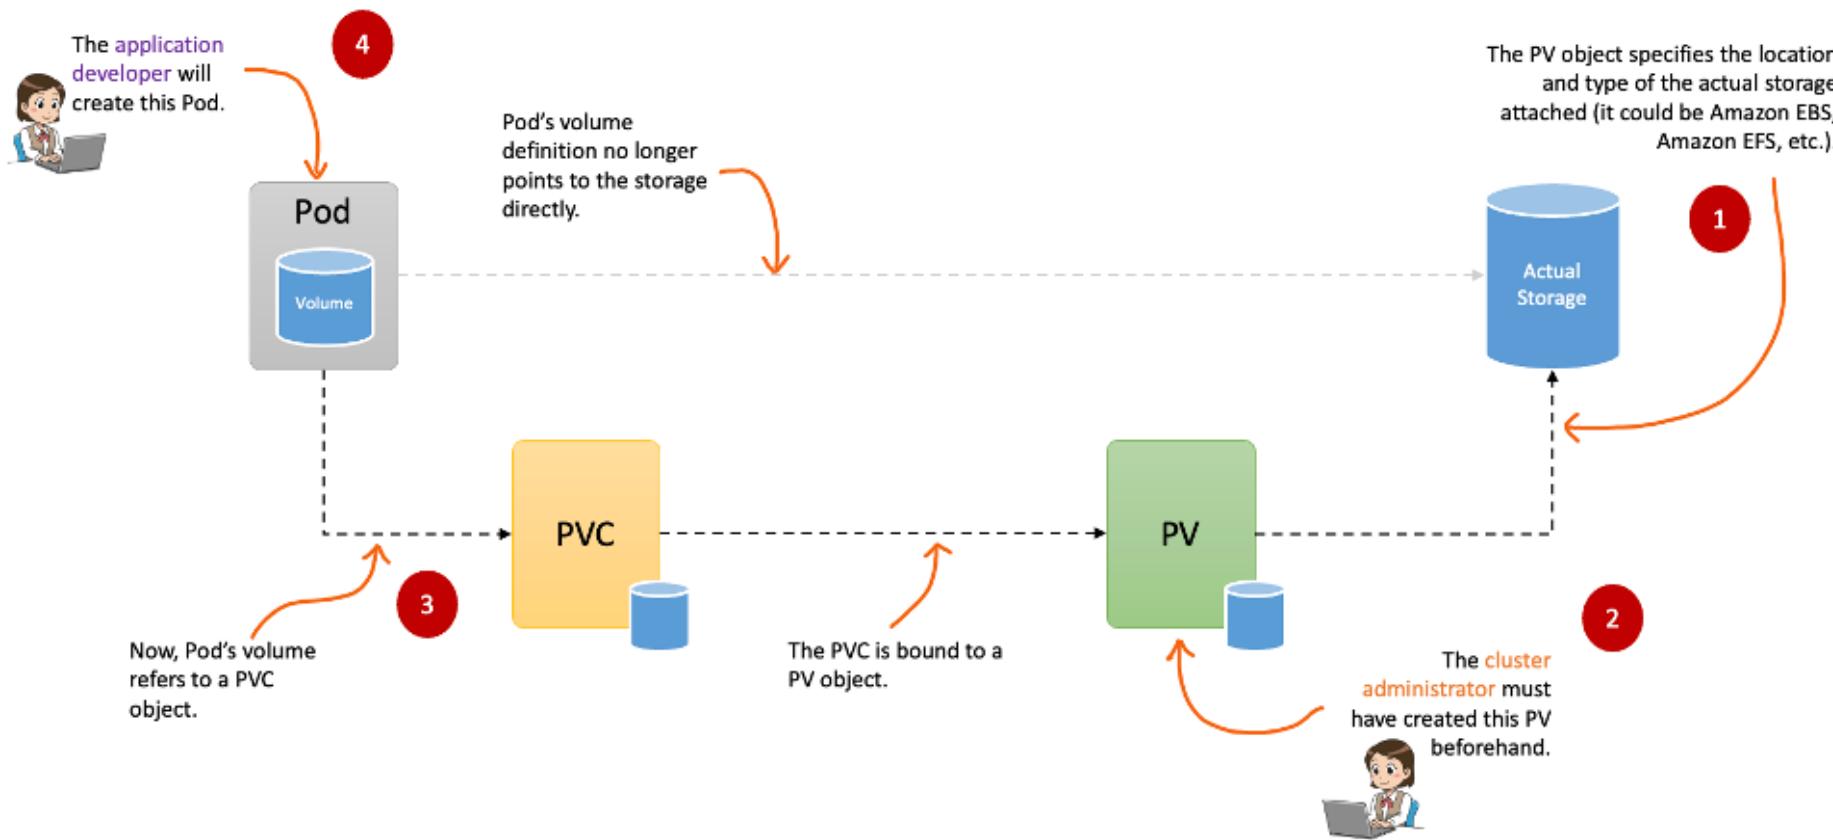
The main goal of storage classes is to eliminate the need for cluster administrators to pre-provision storage and allow them to be created on-demand

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: slow
provisioner: kubernetes.io/gce-pd
parameters:
  type: pd-standard

apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: fast
provisioner: kubernetes.io/gce-pd
parameters:
  type: pd-ssd

apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: ssd-local-path
provisioner: cluster.local/local-path-provisioner
parameters:
  nodePath: /data/ssd
```

PV / MANUAL PROVISIONING

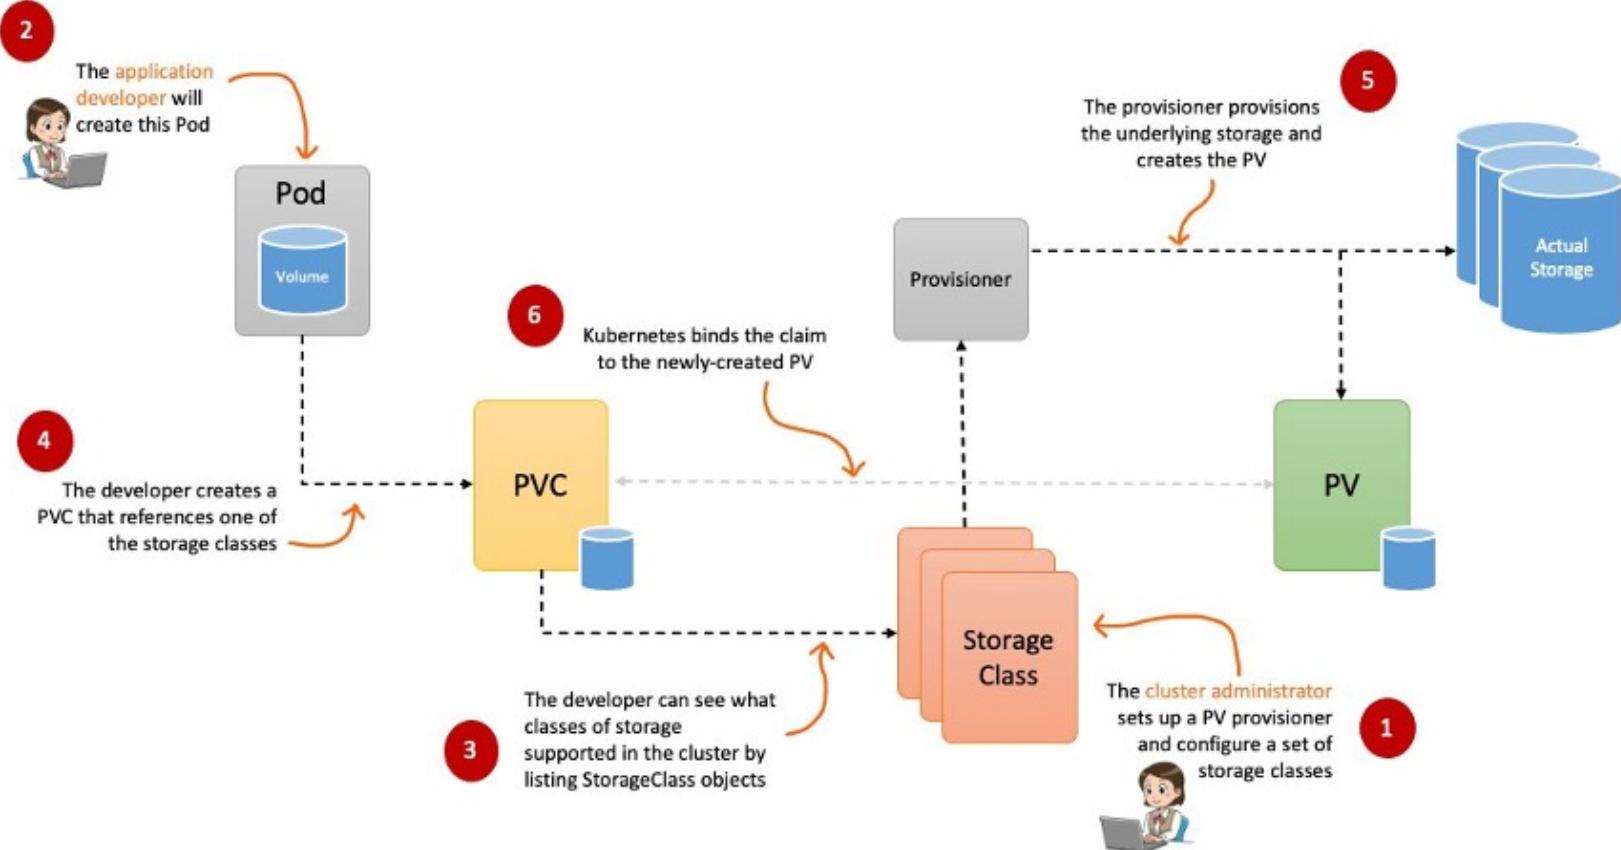


```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-volume01
spec:
  capacity:
    storage: 50Mi
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
  hostPath:
    path: /var/local/data
  ...

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 50Mi
  ...

apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  containers:
    - name: nginx-webserver
      image: nginx
      volumeMounts:
        - name: webserver-vol01
          mountPath: "/usr/share/nginx/html"
  volumes:
    - name: webserver-vol01
      persistentVolumeClaim:
        claimName: myclaim
```

PV / DYNAMIC PROVISIONING



```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: ebs-claim-01
spec:
  storageClassName: gp2-encrypted
  accessModes:
    - ReadWriteOnce
resources:
  requests:
    storage: 10Gi
---
apiVersion: v1
kind: Pod
metadata:
  name: webserver
spec:
  containers:
    - name: nginx-container
      image: nginx:latest
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: test-volume
  volumes:
    - name: test-volume
      persistentVolumeClaim:
        claimName: ebs-claim-01
```

PV / LOCAL-PATH-PROVISIONER

<https://github.com/rancher/local-path-provisioner>

Install

```
kubectl apply -f  
https://raw.githubusercontent.com/rancher/local-path-provisioner/v0.0.24/deploy/local-path-storage.yaml
```

```
// Set local-path as the default dynamic provisioner  
// and the default volume type as 'local'
```

```
kubectl patch storageclass local-path -p '{"metadata":  
{"annotations":{"storageclass.kubernetes.io/is-default-class":"true","defaultVolumeType":"local"}}}'
```

```
kind: ConfigMap  
apiVersion: v1  
metadata:  
  name: local-path-config  
  namespace: local-path-storage  
data:  
  config.json: |-  
    {  
      "nodePathMap": [  
        {  
          "node": "DEFAULT_PATH_FOR_NON_LISTED_NODES",  
          "paths": ["/bbdata"]  
        },  
        {  
          "node": "yasker-lp-dev1",  
          "paths": ["/opt/local-path-provisioner", "/data1"]  
        },  
        {  
          "node": "yasker-lp-dev3",  
          "paths": []  
        }  
      ]  
    }
```

PV / LOCAL-PATH-PROVISIONER

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim2
spec:
  storageClassName: "local-path"
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 50Mi
---
apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  containers:
    - name: nginx-webserver
      image: nginx
      volumeMounts:
        - name: webserver-vol01
          mountPath: "/usr/share/nginx/html"
  volumes:
    - name: webserver-vol01
      persistentVolumeClaim:
        claimName: myclaim2
```

LAB: VOLUMES

Knowing that Postgres stores its data in this folder:

/var/lib/pgsql/data

modify your Postgres pod so it stores Postgres data in a PersistentVolume (manually provisioned) so if the pod is deleted its data will persist

Using PgAdmin or DBBeaver probe the Postgres data was not lost after the pod was previously deleted

<https://github.com/shiriuki/k8s-training/blob/main/kubectl-commands.yaml>

<https://github.com/shiriuki/k8s-training/blob/main/secret.yml>

CONTROLLERS / DEPLOYMENT

Provide a way to define, update, and maintain desired replica counts of Pods

When deployed for the first time, after every update, or restart, behind the scenes it creates a new ReplicaSet k8s resource using the configuration specified in the template section and replicas value

```
// Changes the number of replicas in a deployment
k scale deploy nginx-deployment --replicas 5

// Creates a new ReplicaSet and migrate pods to it
k rollout restart deployment/nginx-deployment

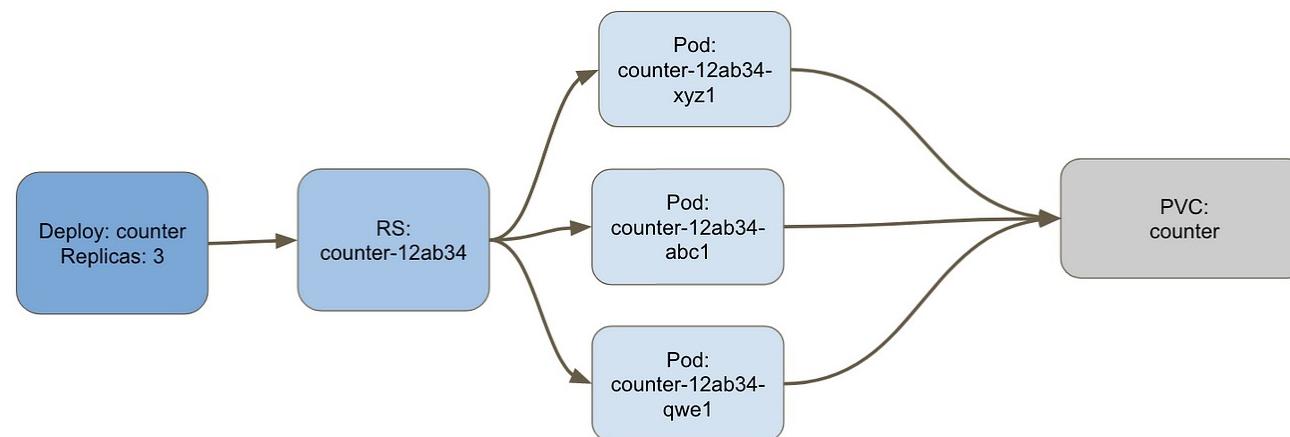
// Go back to the previous ReplicaSet
k rollout undo deployment/nginx-deployment
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx-container
          image: nginx
```

CONTROLLERS / DEPLOYMENT

They should not be used if your deployment directly writes to a persistence volume, as if you have more than 1 replica, that could create inconsistencies because them will share the same PersistentVolume

Persistence in Deployments



<https://medium.com/stakater/k8s-deployments-vs-statefulsets-vs-daemonsets-60582f0c62d4>

CONTROLLERS / STATEFULSET

Manages the deployment and scaling of a set of Pods, and provides guarantee about the ordering and uniqueness of these Pods

It doesn't create ReplicaSet rather itself creates the Pod with a unique naming convention. If you create a StatefulSet with name **counter**, it will create a pod with name **counter-0**, and for multiple replicas of a statefulset, their names will increment like **counter-0, counter-1, counter-2, etc**

```
// Changes the number of replicas in a statefulset
k scale sts/counter --replicas 5

// Creates restarts sts pods
k rollout restart sts/counter

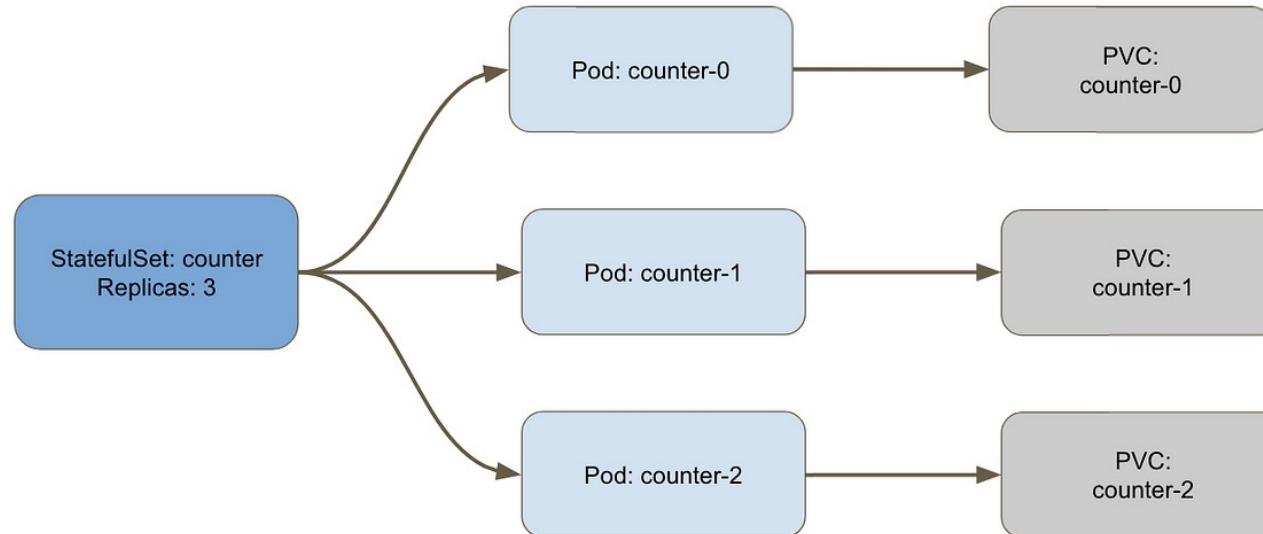
// Go back to the previous configuration
k rollout undo sts/counter
```

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: counter
spec:
  replicas: 1
  selector:
    matchLabels:
      app: counter
  template:
    metadata:
      labels:
        app: counter
    spec:
      containers:
        - name: counter
          image: kahootali/counter:1.1
          volumeMounts:
            - name: counter
              mountPath: /app/
  volumeClaimTemplates:
    - metadata:
        name: counter
      spec:
        accessModes:
          - ReadWriteMany
        resources:
          requests:
            storage: 50Mi
        storageClassName: standard
```

CONTROLLERS / STATEFULSET

Every replica of a stateful set will have its own state, and each of the pods will be creating its own PVC. So a statefulset with 3 replicas will create 3 pods, each having its own Volume, so total 3 PVCs.

Persistence in Statefulsets



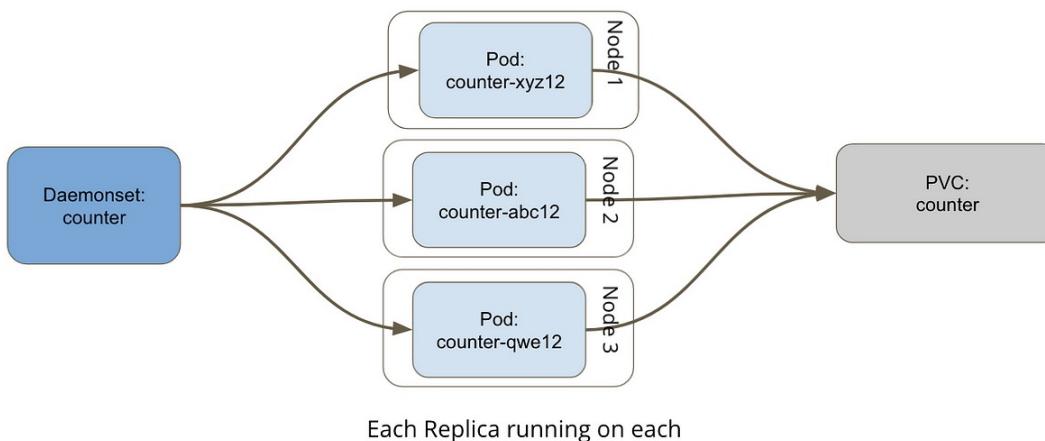
CONTROLLERS / DAEMONSET

Is a controller that ensures that the pod runs on all the nodes of the cluster. If a node is added/removed from a cluster, DaemonSet automatically adds/deletes the pod.

Typical use cases:

- Monitor exporters
- Log collectors

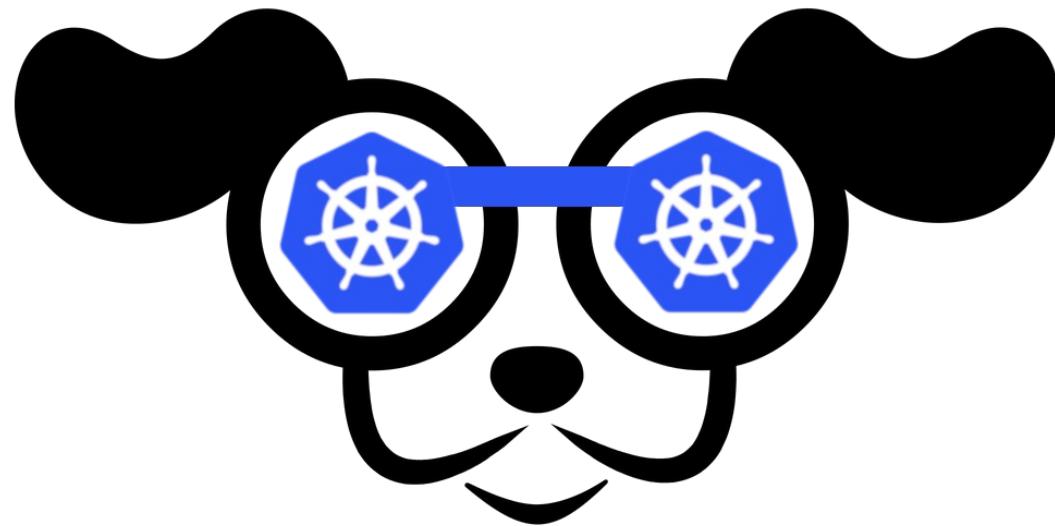
Persistence in Daemonsets



```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluent-bit
  namespace: monitoring
  labels:
    k8s-app: fluent-bit-logging
spec:
  selector:
    matchLabels:
      k8s-app: fluent-bit-logging
  template:
    metadata:
      labels:
        k8s-app: fluent-bit-logging
    spec:
      containers:
        - name: fluent-bit
          image: fluent/fluent-bit:2.1.7
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 2020
          volumeMounts:
            - name: varlog
              mountPath: /var/log
            - name: varlibdockercontainers
              mountPath: /var/lib/docker/containers
              readOnly: true
            - name: fluent-bit-config
              mountPath: /fluent-bit/etc/
      terminationGracePeriodSeconds: 10
  volumes:
    - name: varlog
      hostPath:
        path: /var/log
    - name: varlibdockercontainers
      hostPath:
        path: /var/lib/docker/containers
    - name: fluent-bit-config
      configMap:
        name: fluent-bit-config
```

TOOLS / K9S

<https://k9scli.io/topics/install/>



LAB: FINAL LAB

<https://medium.com/@shubhangi.thakur4532/how-to-deploy-mysql-and-wordpress-on-kubernetes-8ea1260c27dd>

Follow the steps of that article

- Use minikube (no need to use google cloud)
- Use a secret instead of plain text password
- Make mysql db persistent (hint: see this <https://kubernetes.io/docs/tutorials/stateful-application/mysql-wordpress-persistent-volume/>)

Did it work for you? Lets discuss what you did in it

ADDITIONAL TOPICS

Helm charts
Autoscaling (hpa)
Resources / limits
Affinities (pod / node)
QoS / Priorities
Jobs / Cronjobs
Probes
SOPS
Telepresence
Backup (Velero)
Cordon / Drain

Deployment strategies (rolling, blue-green, canary)
Ingress controller, API Gateway
RBAC & Service Accounts
Helmfile
Metrics / Monitoring (Prometheus / Grafana)
Autoscaling (vpa, keda)
Log aggregation (Graylog)
Taints & Tolerations
Service mesh (Istio)

<http://medium.com>

THANK YOU

I hope this training helped you to get start with K8s
and was fun at the same time